

UIT2504 Artificial Intelligence

Local Search Strategies

C. Aravindan
<AravindanC@ssn.edu.in>

Professor of Information Technology
SSN College of Engineering

August 26, 2024

Evaluating a state

- Sometimes, it may be possible to design an evaluation function $f(s)$ that evaluates the “badness” (to be minimized) or “goodness” (to be maximized) of a state s
- In such cases, the most desirable state may be chosen from the working set
- Working set is maintained as a priority queue based on the evaluation function f
- Obviously, the quality of search depends on the evaluation function f

Heuristics


- Usually, such an evaluation function $f(s)$ is designed based on some heuristics $h(s)$ — estimation of cost of reaching a goal state from state s
- For example, can you think of a heuristics for the route finding problem in a map? — Straight line distance (SLD) from the current city to the destination city
- Heuristics should be an easy function to compute!
- $h(s^*)$ should be 0 for any goal state s^*

Admissible Heuristics


- A heuristic h is **admissible** if for every node n , $h(n) \leq h^*(n)$, where $h^*(n)$ is the **true** cost to reach the goal state from n .
- In other words, $f(n) = g(n) + h(n) \leq C^*$, where C^* is the optimal path cost
- An admissible heuristic **never overestimates** the cost to reach the goal, i.e., it is always **optimistic**
- The SLD heuristics and the two heuristics for sliding puzzle problem are examples of admissible heuristics

Consistent Heuristics

- A heuristics is **consistent** if for every node n , $h(n) \leq c(n, a, n') + h(n')$, where n' is a successor of n generated by some action a



Consistent Heuristics



- When h is consistent, we can infer the following

$$\begin{aligned} f(n') &= g(n') + h(n') \\ &= g(n) + c(n, a, n') + h(n') \\ &\geq g(n) + h(n) \\ &\geq f(n) \end{aligned}$$

- That means, evaluation function f is monotonic — it is non-decreasing along any path
- Every consistent heuristics is also admissible

Questions?

Iterative improvement algorithms

- In several problems, the path is irrelevant and a goal-state is a solution we are looking for

Iterative improvement algorithms

- In several problems, the path is irrelevant and a goal-state is a solution we are looking for
- With state space = set of “complete” configurations,
 - Find an optimal configuration (eg. TSP, maximal matching in a bipartite graph, “weights” that minimize error on the examples)

Iterative improvement algorithms

- In several problems, the path is irrelevant and a goal-state is a solution we are looking for
- With state space = set of “complete” configurations,
 - Find an optimal configuration (eg. TSP, maximal matching in a bipartite graph, “weights” that minimize error on the examples)
 - Find a configuration that satisfies some constraints (eg. Timetable generation, n -queens problem, stable matching)

Iterative improvement algorithms


- In several problems, the path is irrelevant and a goal-state is a solution we are looking for
- With state space = set of “complete” configurations,
 - Find an optimal configuration (eg. TSP, maximal matching in a bipartite graph, “weights” that minimize error on the examples)
 - Find a configuration that satisfies some constraints (eg. Timetable generation, n -queens problem, stable matching)
- In such cases, we can use **iterative improvement algorithms** — “keep a single current state and try to improve it”

Example: TSP

- “Complete formulation” — any Hamiltonian circuit is a state


Example: TSP

- “Complete formulation” — any Hamiltonian circuit is a state
- Start with any complete state (any Hamiltonian circuit)
- Action: Pairwise exchanges — replace edges (u, v) and (u', v') with edges (u, v') and (u', v)



Example: TSP

- “Complete formulation” — any Hamiltonian circuit is a state
- Start with any complete state (any Hamiltonian circuit)
- Action: Pairwise exchanges — replace edges (u, v) and (u', v') with edges (u, v') and (u', v)



- Able to get within 1% of optimal solution very quickly, even with thousands of cities (good, as an **approximation algorithm**)

Example: n -queens problem


- “Complete formulation” — n queens on board, one per column

Example: n -queens problem

- “Complete formulation” — n queens on board, one per column
- Start with any state — any placement of n queens, one per column, on a chess board


Example: n -queens problem

- “Complete formulation” — n queens on board, one per column
- Start with any state — any placement of n queens, one per column, on a chess board
- Action: Move a queen, within its column, to reduce the number of conflicts



Example: n -queens problem

- “Complete formulation” — n queens on board, one per column
- Start with any state — any placement of n queens, one per column, on a chess board
- Action: Move a queen, within its column, to reduce the number of conflicts



- Able to solve instantaneously even for very large n

Outline of Hill Climbing Algorithm

```
def hill_climbing(problem):  
    current = problem.initial()  
    while True:  
        neighbor = max( problem.children(current) )  
        if problem.value(neighbor) <=  
            problem.value(current):  
            break  
        current = neighbor  
    return current
```

Example: 8-queens problem


18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♔	13	16	13	16
♔	14	17	15	♔	14	16	16
17	♔	16	18	15	♔	15	♔
18	14	♔	15	15	14	♔	16
14	14	13	17	12	14	12	18

Example: 8-queens problem


18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♔	13	16	13	16
♔	14	17	15	♔	14	16	16
17	♔	16	18	15	♔	15	♔
18	14	♔	15	15	14	♔	16
14	14	13	17	12	14	12	18

- One of the square with $h = 12$ is chosen and the corresponding queen is moved there

Example: 8-queens problem




Example: 8-queens problem




- $h(s) = 1$

Example: 8-queens problem



- $h(s) = 1$
- And, we have hit a local minimum!

Hill Climbing Search



Issues in Greedy Local Search

- Local maxima / minima

Issues in Greedy Local Search

- Local maxima / minima
- Ridges — sequence of local maxima that is very difficult for the greedy algorithm to navigate

Issues in Greedy Local Search

- Local maxima / minima
- Ridges — sequence of local maxima that is very difficult for the greedy algorithm to navigate
- Plateaux — flat local maximum or a shoulder


Issues in Greedy Local Search

- **Local maxima / minima**
- **Ridges** — sequence of local maxima that is very difficult for the greedy algorithm to navigate
- **Plateaux** — flat local maximum or a shoulder

Local maxima is a serious problem

Empirical analysis of 8-queens problem reveals that the greedy hill-climbing algorithm gets stuck 86% of the time

Ridges



Variations of Hill Climbing

- Random sideways moves

Variations of Hill Climbing

- **Random sideways moves** — can escape from a shoulder, but gets trapped in a local maxima

Variations of Hill Climbing

- **Random sideways moves** — can escape from a shoulder, but gets trapped in a local maxima — number of sideways moves may be limited

Variations of Hill Climbing

- **Random sideways moves** — can escape from a shoulder, but gets trapped in a local maxima — number of sideways moves may be limited — number of instances solved for the 8-queens problem increases from 14% to 94%

Variations of Hill Climbing

- **Random sideways moves** — can escape from a shoulder, but gets trapped in a local maxima — number of sideways moves may be limited — number of instances solved for the 8-queens problem increases from 14% to 94%
- **Stochastic Hill Climbing**

Variations of Hill Climbing

- **Random sideways moves** — can escape from a shoulder, but gets trapped in a local maxima — number of sideways moves may be limited — number of instances solved for the 8-queens problem increases from 14% to 94%
- **Stochastic Hill Climbing** — chooses at random, among all uphill moves

Variations of Hill Climbing

- **Random sideways moves** — can escape from a shoulder, but gets trapped in a local maxima — number of sideways moves may be limited — number of instances solved for the 8-queens problem increases from 14% to 94%
- **Stochastic Hill Climbing** — chooses at random, among all uphill moves — probability of selection can depend on the steepness of the ascent

Variations of Hill Climbing

- **Random sideways moves** — can escape from a shoulder, but gets trapped in a local maxima — number of sideways moves may be limited — number of instances solved for the 8-queens problem increases from 14% to 94%
- **Stochastic Hill Climbing** — chooses at random, among all uphill moves — probability of selection can depend on the steepness of the ascent — example of a **Randomized algorithm**

Variations of Hill Climbing

- **Random sideways moves** — can escape from a shoulder, but gets trapped in a local maxima — number of sideways moves may be limited — number of instances solved for the 8-queens problem increases from 14% to 94%
- **Stochastic Hill Climbing** — chooses at random, among all uphill moves — probability of selection can depend on the steepness of the ascent — example of a **Randomized algorithm**
- **First Choice Hill Climbing** — randomly generate the successors, until one better than the current is generated

Variations of Hill Climbing

- **Random sideways moves** — can escape from a shoulder, but gets trapped in a local maxima — number of sideways moves may be limited — number of instances solved for the 8-queens problem increases from 14% to 94%
- **Stochastic Hill Climbing** — chooses at random, among all uphill moves — probability of selection can depend on the steepness of the ascent — example of a **Randomized algorithm**
- **First Choice Hill Climbing** — randomly generate the successors, until one better than the current is generated
- **Random Restart**

Variations of Hill Climbing

- **Random sideways moves** — can escape from a shoulder, but gets trapped in a local maxima — number of sideways moves may be limited — number of instances solved for the 8-queens problem increases from 14% to 94%
- **Stochastic Hill Climbing** — chooses at random, among all uphill moves — probability of selection can depend on the steepness of the ascent — example of a **Randomized algorithm**
- **First Choice Hill Climbing** — randomly generate the successors, until one better than the current is generated
- **Random Restart** — enough restarts may make this algorithm complete

Variations of Hill Climbing

- **Random sideways moves** — can escape from a shoulder, but gets trapped in a local maxima — number of sideways moves may be limited — number of instances solved for the 8-queens problem increases from 14% to 94%
- **Stochastic Hill Climbing** — chooses at random, among all uphill moves — probability of selection can depend on the steepness of the ascent — example of a **Randomized algorithm**
- **First Choice Hill Climbing** — randomly generate the successors, until one better than the current is generated
- **Random Restart** — enough restarts may make this algorithm complete — if each hill-climbing has a probability p of success, then $1/p$ restarts are expected

Variations of Hill Climbing

- **Random sideways moves** — can escape from a shoulder, but gets trapped in a local maxima — number of sideways moves may be limited — number of instances solved for the 8-queens problem increases from 14% to 94%
- **Stochastic Hill Climbing** — chooses at random, among all uphill moves — probability of selection can depend on the steepness of the ascent — example of a **Randomized algorithm**
- **First Choice Hill Climbing** — randomly generate the successors, until one better than the current is generated
- **Random Restart** — enough restarts may make this algorithm complete — if each hill-climbing has a probability p of success, then $1/p$ restarts are expected — for 8-queens, $p \approx 0.14$, and so roughly 7 restarts are expected

Questions?

Simulated Annealing

- One interesting variation of hill climbing is to adopt the concept of **simulated annealing**

Simulated Annealing

- One interesting variation of hill climbing is to adopt the concept of **simulated annealing**
- For example, consider a ball set to roll on a state landscape

Simulated Annealing

- One interesting variation of hill climbing is to adopt the concept of **simulated annealing**
- For example, consider a ball set to roll on a state landscape
- The ball simply follows the rules of gravity and moves towards nearby valley

Simulated Annealing

- One interesting variation of hill climbing is to adopt the concept of **simulated annealing**
- For example, consider a ball set to roll on a state landscape
- The ball simply follows the rules of gravity and moves towards nearby valley
- The ball needs to make some uphill moves to escape from local minima!

Simulated Annealing

- One interesting variation of hill climbing is to adopt the concept of **simulated annealing**
- For example, consider a ball set to roll on a state landscape
- The ball simply follows the rules of gravity and moves towards nearby valley
- The ball needs to make some uphill moves to escape from local minima!
- Imagine applying just enough force for it to escape from all local minima but not from global minima

Simulated Annealing

- One interesting variation of hill climbing is to adopt the concept of **simulated annealing**
- For example, consider a ball set to roll on a state landscape
- The ball simply follows the rules of gravity and moves towards nearby valley
- The ball needs to make some uphill moves to escape from local minima!
- Imagine applying just enough force for it to escape from all local minima but not from global minima
- How to find that “just enough force”?

Simulated Annealing

- Similar to the metallurgical process of annealing

Simulated Annealing

- Similar to the metallurgical process of annealing
- Start with a high “temperature”

Simulated Annealing

- Similar to the metallurgical process of annealing
- Start with a high “temperature” — probability of selecting a bad move is high

Simulated Annealing

- Similar to the metallurgical process of annealing
- Start with a high “temperature” — probability of selecting a bad move is high
- Slowly reduce the “temperature”

Simulated Annealing

- Similar to the metallurgical process of annealing
- Start with a high “temperature” — probability of selecting a bad move is high
- Slowly reduce the “temperature” — probability of selecting a bad move reduces slowly

Simulated Annealing

- Similar to the metallurgical process of annealing
- Start with a high “temperature” — probability of selecting a bad move is high
- Slowly reduce the “temperature” — probability of selecting a bad move reduces slowly
- “Schedule” of reducing the temperature is very critical

Simulated Annealing

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  current  $\leftarrow$  problem.INITIAL
  for  $t = 1$  to  $\infty$  do
     $T \leftarrow$  schedule( $t$ )
    if  $T = 0$  then return current
    next  $\leftarrow$  a randomly selected successor of current
     $\Delta E \leftarrow$  VALUE(current) - VALUE(next)
    if  $\Delta E > 0$  then current  $\leftarrow$  next
    else current  $\leftarrow$  next only with probability  $e^{-\Delta E/T}$ 
```

Questions?