

UIT2504 Artificial Intelligence

Playing Games

C. Aravindan
<AravindanC@ssn.edu.in>

Professor of Information Technology
SSN College of Engineering

September 04, 2024




Problem Solving

- Problem formulation in state space
- Basic Search Strategies
- Heuristics — best-first search
- A^* Search
- Local Search Strategies — Hill climbing and its variants, simulated annealing, genetic algorithms
- Searching in continuous domains

Problem Solving

- Problem formulation in state space
- Basic Search Strategies
- Heuristics — best-first search
- A^* Search
- Local Search Strategies — Hill climbing and its variants, simulated annealing, genetic algorithms
- Searching in continuous domains
- Will these ideas be useful for taking decisions in multi-player games?

Search Strategies for Playing Games



- Search tree is huge! — in chess, average branching factor is 35 and a game may go to a depth of 100 (50 moves by each player) — 35^{100}

- Search tree is huge! — in chess, average branching factor is 35 and a game may go to a depth of 100 (50 moves by each player) — 35^{100}
- There may be three approaches to handle multi-agent environments:

- Search tree is huge! — in chess, average branching factor is 35 and a game may go to a depth of 100 (50 moves by each player) — 35^{100}
- There may be three approaches to handle multi-agent environments:
 - Consider all agents in aggregate as an **economy**

- Search tree is huge! — in chess, average branching factor is 35 and a game may go to a depth of 100 (50 moves by each player) — 35^{100}
- There may be three approaches to handle multi-agent environments:
 - Consider all agents in aggregate as an **economy**
 - Consider all the adversarial agents as parts of the environment — that makes the environment nondeterministic

- Search tree is huge! — in chess, average branching factor is 35 and a game may go to a depth of 100 (50 moves by each player) — 35^{100}
- There may be three approaches to handle multi-agent environments:
 - Consider all agents in aggregate as an **economy**
 - Consider all the adversarial agents as parts of the environment — that makes the environment nondeterministic
 - Explicitly model the adversarial agents and develop special techniques for game-tree search

- Search tree is huge! — in chess, average branching factor is 35 and a game may go to a depth of 100 (50 moves by each player) — 35^{100}
- There may be three approaches to handle multi-agent environments:
 - Consider all agents in aggregate as an **economy**
 - Consider all the adversarial agents as parts of the environment — that makes the environment nondeterministic
 - Explicitly model the adversarial agents and develop special techniques for game-tree search
- We will consider two-player **Zero-sum games of perfect information**

- Search tree is huge! — in chess, average branching factor is 35 and a game may go to a depth of 100 (50 moves by each player) — 35^{100}
- There may be three approaches to handle multi-agent environments:
 - Consider all agents in aggregate as an **economy**
 - Consider all the adversarial agents as parts of the environment — that makes the environment nondeterministic
 - Explicitly model the adversarial agents and develop special techniques for game-tree search
- We will consider two-player **Zero-sum games of perfect information**
- We will call our two players *MAX* and *MIN*

Game Formulation

- States and the initial state S_0

Game Formulation

- States and the initial state S_0
- **TO_MOVE(s)**: The player whose turn it is to move in state s

Game Formulation

- States and the initial state S_0
- $TO_MOVE(s)$: The player whose turn it is to move in state s
- $ACTIONS(s)$: set of legal moves in a state s

Game Formulation

- States and the initial state S_0
- $TO_MOVE(s)$: The player whose turn it is to move in state s
- $ACTIONS(s)$: set of legal moves in a state s
- $RESULT(s,a)$: the transition model that defines the result of a move a in a state s


Game Formulation

- States and the initial state S_0
- **TO_MOVE(s)**: The player whose turn it is to move in state s
- **ACTIONS(s)**: set of legal moves in a state s
- **RESULT(s,a)**: the transition model that defines the result of a move a in a state s
- **IS_TERMINAL(s)**: Terminal test — 'true' when the game is over — terminal states

Game Formulation

- States and the initial state S_0
- $TO_MOVE(s)$: The player whose turn it is to move in state s
- $ACTIONS(s)$: set of legal moves in a state s
- $RESULT(s,a)$: the transition model that defines the result of a move a in a state s
- $IS_TERMINAL(s)$: Terminal test — 'true' when the game is over — terminal states
- $UTILITY(s,p)$: Utility function — an objective function that defines the final numeric value to a player p when the game ends in a terminal state s — in chess, outcome is a win, loss, or draw, with values $+1, 0, \frac{1}{2}$

Optimal Decisions



Minimax Decision

- Optimal strategy for deterministic games
- Utility values percolate up from terminal states
- At MAX level, choose successor with highest utility
- At MIN level, choose successor with lowest utility (note that utility is from MAX's point of view)

Minimax Decision

- Optimal strategy for deterministic games
- Utility values percolate up from terminal states
- At MAX level, choose successor with highest utility
- At MIN level, choose successor with lowest utility (note that utility is from MAX's point of view)

MINIMAX(s) =

$$\begin{cases} \text{Utility}(s, \text{MAX}) & \text{if IS_TERMINAL}(s) \\ \max_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if TO_MOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if TO_MOVE}(s) = \text{MIN} \end{cases}$$

Minimax Algorithm

function MINIMAX-SEARCH(*game*, *state*) **returns** an action

player \leftarrow *game*.TO-MOVE(*state*)

value, *move* \leftarrow MAX-VALUE(*game*, *state*)

return *move*

function MAX-VALUE(*game*, *state*) **returns** a (*utility*, *move*) pair

if *game*.IS-TERMINAL(*state*) **then return** *game*.UTILITY(*state*, *player*), null

v $\leftarrow -\infty$

for each *a* **in** *game*.ACTIONS(*state*) **do**

v2, *a2* \leftarrow MIN-VALUE(*game*, *game*.RESULT(*state*, *a*))

if *v2* > *v* **then**

v, *move* \leftarrow *v2*, *a*

return *v*, *move*

function MIN-VALUE(*game*, *state*) **returns** a (*utility*, *move*) pair

if *game*.IS-TERMINAL(*state*) **then return** *game*.UTILITY(*state*, *player*), null

v $\leftarrow +\infty$

for each *a* **in** *game*.ACTIONS(*state*) **do**


v2, *a2* \leftarrow MAX-VALUE(*game*, *game*.RESULT(*state*, *a*))

if *v2* < *v* **then**

v, *move* \leftarrow *v2*, *a*

return *v*, *move*

Optimal Decisions



- Complete?

- Complete? — Yes, if the tree is finite

- Complete? — Yes, if the tree is finite — complete depth-first search

- Complete? — Yes, if the tree is finite — complete depth-first search
- Optimal?

- Complete? — Yes, if the tree is finite — complete depth-first search
- Optimal? — Yes, if the opponent is optimal

- Complete? — Yes, if the tree is finite — complete depth-first search
- Optimal? — Yes, if the opponent is optimal
- Space Complexity?

- Complete? — Yes, if the tree is finite — complete depth-first search
- Optimal? — Yes, if the opponent is optimal
- Space Complexity? — $O(bm)$

- Complete? — Yes, if the tree is finite — complete depth-first search
- Optimal? — Yes, if the opponent is optimal
- Space Complexity? — $O(bm)$ — may be reduced to $O(m)$ if successors are generated one at a time


- Complete? — Yes, if the tree is finite — complete depth-first search
- Optimal? — Yes, if the opponent is optimal
- Space Complexity? — $O(bm)$ — may be reduced to $O(m)$ if successors are generated one at a time
- Time Complexity?

- Complete? — Yes, if the tree is finite — complete depth-first search
- Optimal? — Yes, if the opponent is optimal
- Space Complexity? — $O(bm)$ — may be reduced to $O(m)$ if successors are generated one at a time
- Time Complexity? — $O(b^m)$

- Complete? — Yes, if the tree is finite — complete depth-first search
- Optimal? — Yes, if the opponent is optimal
- Space Complexity? — $O(bm)$ — may be reduced to $O(m)$ if successors are generated one at a time
- Time Complexity? — $O(b^m)$
- Impractical for non-trivial games such as chess — 35^{100}


Multiplayer Games

to move
A




Questions?


Pruning




Pruning




Pruning




Pruning



Pruning



Pruning




Alpha-Beta Pruning

- α : Value of the best choice we have found so far along a path for MAX — α = “at least” — lower bound for MAX
- β : Value of the best choice we have found so far along a path for MIN — β = “at most” — upper bound for MAX

Alpha-Beta Pruning

- α : Value of the best choice we have found so far along a path for MAX — α = “at least” — lower bound for MAX
- β : Value of the best choice we have found so far along a path for MIN — β = “at most” — upper bound for MAX



Alpha-Beta Search Algorithm

function ALPHA-BETA-SEARCH(*game*, *state*) **returns** an action

$\text{player} \leftarrow \text{game.TO-MOVE}(\text{state})$

$\text{value}, \text{move} \leftarrow \text{MAX-VALUE}(\text{game}, \text{state}, -\infty, +\infty)$

return *move*

function MAX-VALUE(*game*, *state*, α , β) **returns** a (*utility*, *move*) pair

if *game.IS-TERMINAL*(*state*) **then return** *game.UTILITY*(*state*, *player*), *null*

$v \leftarrow -\infty$

for each *a* **in** *game.ACTIONS*(*state*) **do**

$v2, a2 \leftarrow \text{MIN-VALUE}(\text{game}, \text{game.RESULT}(\text{state}, a), \alpha, \beta)$

if $v2 > v$ **then**

$v, \text{move} \leftarrow v2, a$

$\alpha \leftarrow \text{MAX}(\alpha, v)$

if $v \geq \beta$ **then return** *v*, *move*

return *v*, *move*

function MIN-VALUE(*game*, *state*, α , β) **returns** a (*utility*, *move*) pair

if *game.IS-TERMINAL*(*state*) **then return** *game.UTILITY*(*state*, *player*), *null*

$v \leftarrow +\infty$

for each *a* **in** *game.ACTIONS*(*state*) **do**

$v2, a2 \leftarrow \text{MAX-VALUE}(\text{game}, \text{game.RESULT}(\text{state}, a), \alpha, \beta)$

if $v2 < v$ **then**

$v, \text{move} \leftarrow v2, a$

$\beta \leftarrow \text{MIN}(\beta, v)$

if $v \leq \alpha$ **then return** *v*, *move*

return *v*, *move*

- α - β pruning does not affect the completeness or optimality of the minimax algorithm

Complexities


- α - β pruning does not affect the completeness or optimality of the minimax algorithm
- Space complexity is also same as that of the minimax algorithm

Complexities

- α - β pruning does not affect the completeness or optimality of the minimax algorithm
- Space complexity is also same as that of the minimax algorithm
- If the successors can be generated in best-first order (can it be done?), number of nodes to be examined may be halved

Complexities

- α - β pruning does not affect the completeness or optimality of the minimax algorithm
- Space complexity is also same as that of the minimax algorithm
- If the successors can be generated in best-first order (can it be done?), number of nodes to be examined may be halved



Move Ordering

- If best move ordering can be achieved, time taken may be halved

Move Ordering

- If best move ordering can be achieved, time taken may be halved — $O(b^{m/2})$

Move Ordering

- If best move ordering can be achieved, time taken may be halved — $O(b^{m/2})$ — which is same as $O\left((\sqrt{b})^m\right)$ —

Move Ordering

- If best move ordering can be achieved, time taken may be halved — $O(b^{m/2})$ — which is same as $O\left(\left(\sqrt{b}\right)^m\right)$ — effective branching factor is reduced to \sqrt{b}

Move Ordering

- If best move ordering can be achieved, time taken may be halved — $O(b^{m/2})$ — which is same as $O\left((\sqrt{b})^m\right)$ — effective branching factor is reduced to \sqrt{b}
- If the successors are examined in a random order, time complexity will be roughly $O(b^{3m/4})$

Move Ordering

- If best move ordering can be achieved, time taken may be halved — $O(b^{m/2})$ — which is same as $O\left((\sqrt{b})^m\right)$ — effective branching factor is reduced to \sqrt{b}
- If the successors are examined in a random order, time complexity will be roughly $O(b^{3m/4})$
- Heuristics may be used to order the moves

Move Ordering

- If best move ordering can be achieved, time taken may be halved — $O(b^{m/2})$ — which is same as $O\left((\sqrt{b})^m\right)$ — effective branching factor is reduced to \sqrt{b}
- If the successors are examined in a random order, time complexity will be roughly $O(b^{3m/4})$
- Heuristics may be used to order the moves — in chess, captures first, then threats, then forward moves, and then backward moves

Move Ordering

- If best move ordering can be achieved, time taken may be halved — $O(b^{m/2})$ — which is same as $O\left((\sqrt{b})^m\right)$ — effective branching factor is reduced to \sqrt{b}
- If the successors are examined in a random order, time complexity will be roughly $O(b^{3m/4})$
- Heuristics may be used to order the moves — in chess, captures first, then threats, then forward moves, and then backward moves
- **Dynamic move ordering** — moves found to be best in the past (learning) — or, iterative-deepening

Move Ordering

- If best move ordering can be achieved, time taken may be halved — $O(b^{m/2})$ — which is same as $O\left((\sqrt{b})^m\right)$ — effective branching factor is reduced to \sqrt{b}
- If the successors are examined in a random order, time complexity will be roughly $O(b^{3m/4})$
- Heuristics may be used to order the moves — in chess, captures first, then threats, then forward moves, and then backward moves
- **Dynamic move ordering** — moves found to be best in the past (learning) — or, iterative-deepening
- **Transposition table** — hash the α - β values for future use

Type A and Type B Strategies

- Even with alpha-beta pruning and other techniques such as move ordering, minimax algorithm may not work for games such as Chess and Go

Type A and Type B Strategies

- Even with alpha-beta pruning and other techniques such as move ordering, minimax algorithm may not work for games such as Chess and Go
- Claude Shannon proposed two additional strategies:
- **Type A strategy**: consider all possible moves to certain depth and then use a heuristic evaluation function to estimate the utilities of the states at that depth (wide but shallow strategy)

Type A and Type B Strategies

- Even with alpha-beta pruning and other techniques such as move ordering, minimax algorithm may not work for games such as Chess and Go
- Claude Shannon proposed two additional strategies:
- **Type A strategy**: consider all possible moves to certain depth and then use a heuristic evaluation function to estimate the utilities of the states at that depth (wide but shallow strategy)
- **Type B Strategy**: ignore moves that look bad, and follow promising paths as far as possible (deep but narrow strategy)

Questions?