

# UIT2504 Artificial Intelligence

## Imperfect Decisions in Games

C. Aravindan  
<AravindanC@ssn.edu.in>

Professor of Information Technology  
SSN College of Engineering

September 11, 2024

# Game Formulation

- States and the initial state  $S_0$

# Game Formulation

- States and the initial state  $S_0$
- **TO\_MOVE(s)**: The player whose turn it is to move in state  $s$

# Game Formulation

- States and the initial state  $S_0$
- $TO\_MOVE(s)$ : The player whose turn it is to move in state  $s$
- $ACTIONS(s)$ : set of legal moves in a state  $s$

# Game Formulation

- States and the initial state  $S_0$
- $TO\_MOVE(s)$ : The player whose turn it is to move in state  $s$
- $ACTIONS(s)$ : set of legal moves in a state  $s$
- $RESULT(s,a)$ : the transition model that defines the result of a move  $a$  in a state  $s$


# Game Formulation

- States and the initial state  $S_0$
- **TO\_MOVE(s)**: The player whose turn it is to move in state  $s$
- **ACTIONS(s)**: set of legal moves in a state  $s$
- **RESULT(s,a)**: the transition model that defines the result of a move  $a$  in a state  $s$
- **IS\_TERMINAL(s)**: Terminal test — 'true' when the game is over — terminal states

# Game Formulation

- States and the initial state  $S_0$
- $TO\_MOVE(s)$ : The player whose turn it is to move in state  $s$
- $ACTIONS(s)$ : set of legal moves in a state  $s$
- $RESULT(s,a)$ : the transition model that defines the result of a move  $a$  in a state  $s$
- $IS\_TERMINAL(s)$ : Terminal test — 'true' when the game is over — terminal states
- $UTILITY(s,p)$ : Utility function — an objective function that defines the final numeric value to a player  $p$  when the game ends in a terminal state  $s$  — in chess, outcome is a win, loss, or draw, with values  $+1, 0, \frac{1}{2}$

# Search Strategies for Playing Games





# Minimax Decision

- Optimal strategy for deterministic games
- Utility values percolate up from terminal states
- At MAX level, choose successor with highest utility
- At MIN level, choose successor with lowest utility (note that utility is from MAX's point of view)

# Minimax Decision

- Optimal strategy for deterministic games
- Utility values percolate up from terminal states
- At MAX level, choose successor with highest utility
- At MIN level, choose successor with lowest utility (note that utility is from MAX's point of view)

MINIMAX( $s$ ) =

$$\begin{cases} \text{Utility}(s, \text{MAX}) & \text{if IS\_TERMINAL}(s) \\ \max_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if TO\_MOVE}(s) = \text{MAX} \\ \min_{a \in \text{ACTIONS}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if TO\_MOVE}(s) = \text{MIN} \end{cases}$$


- Complete? — Yes, if the tree is finite — complete depth-first search
- Optimal? — Yes, if the opponent is optimal
- Space Complexity? —  $O(bm)$  — may be reduced to  $O(m)$  if successors are generated one at a time
- Time Complexity? —  $O(b^m)$
- Impractical for non-trivial games such as chess —  $35^{100}$

# Alpha-Beta Pruning

- $\alpha$  : Value of the best choice we have found so far along a path for MAX —  $\alpha$  = “at least” — lower bound for MAX
- $\beta$  : Value of the best choice we have found so far along a path for MIN —  $\beta$  = “at most” — upper bound for MAX

# Alpha-Beta Pruning

- $\alpha$  : Value of the best choice we have found so far along a path for MAX —  $\alpha$  = “at least” — lower bound for MAX
- $\beta$  : Value of the best choice we have found so far along a path for MIN —  $\beta$  = “at most” — upper bound for MAX



# Move Ordering

- If best move ordering can be achieved, time taken may be halved

# Move Ordering

- If best move ordering can be achieved, time taken may be halved —  $O(b^{m/2})$

# Move Ordering

- If best move ordering can be achieved, time taken may be halved —  $O(b^{m/2})$  — which is same as  $O\left((\sqrt{b})^m\right)$  —



# Move Ordering

- If best move ordering can be achieved, time taken may be halved —  $O(b^{m/2})$  — which is same as  $O\left(\left(\sqrt{b}\right)^m\right)$  — effective branching factor is reduced to  $\sqrt{b}$

# Move Ordering

- If best move ordering can be achieved, time taken may be halved —  $O(b^{m/2})$  — which is same as  $O\left(\left(\sqrt{b}\right)^m\right)$  — effective branching factor is reduced to  $\sqrt{b}$
- If the successors are examined in a random order, time complexity will be roughly  $O(b^{3m/4})$

# Move Ordering

- If best move ordering can be achieved, time taken may be halved —  $O(b^{m/2})$  — which is same as  $O\left((\sqrt{b})^m\right)$  — effective branching factor is reduced to  $\sqrt{b}$
- If the successors are examined in a random order, time complexity will be roughly  $O(b^{3m/4})$
- Heuristics may be used to order the moves

# Move Ordering

- If best move ordering can be achieved, time taken may be halved —  $O(b^{m/2})$  — which is same as  $O\left((\sqrt{b})^m\right)$  — effective branching factor is reduced to  $\sqrt{b}$
- If the successors are examined in a random order, time complexity will be roughly  $O(b^{3m/4})$
- Heuristics may be used to order the moves — in chess, captures first, then threats, then forward moves, and then backward moves

# Move Ordering

- If best move ordering can be achieved, time taken may be halved —  $O(b^{m/2})$  — which is same as  $O\left((\sqrt{b})^m\right)$  — effective branching factor is reduced to  $\sqrt{b}$
- If the successors are examined in a random order, time complexity will be roughly  $O(b^{3m/4})$
- Heuristics may be used to order the moves — in chess, captures first, then threats, then forward moves, and then backward moves
- **Dynamic move ordering** — moves found to be best in the past (learning) — or, iterative-deepening

# Move Ordering

- If best move ordering can be achieved, time taken may be halved —  $O(b^{m/2})$  — which is same as  $O((\sqrt{b})^m)$  — effective branching factor is reduced to  $\sqrt{b}$
- If the successors are examined in a random order, time complexity will be roughly  $O(b^{3m/4})$
- Heuristics may be used to order the moves — in chess, captures first, then threats, then forward moves, and then backward moves
- **Dynamic move ordering** — moves found to be best in the past (learning) — or, iterative-deepening
- **Transposition table** — hash the  $\alpha$ - $\beta$  values for future use

# Type A and Type B Strategies

- Even with alpha-beta pruning and other techniques such as move ordering, minimax algorithm may not work for games such as Chess and Go

# Type A and Type B Strategies

- Even with alpha-beta pruning and other techniques such as move ordering, minimax algorithm may not work for games such as Chess and Go
- Claude Shannon proposed two additional strategies:
- **Type A strategy**: consider all possible moves to certain depth and then use a heuristic evaluation function to estimate the utilities of the states at that depth (wide but shallow strategy)



# Type A and Type B Strategies

- Even with alpha-beta pruning and other techniques such as move ordering, minimax algorithm may not work for games such as Chess and Go
- Claude Shannon proposed two additional strategies:
- **Type A strategy**: consider all possible moves to certain depth and then use a heuristic evaluation function to estimate the utilities of the states at that depth (wide but shallow strategy)
- **Type B Strategy**: ignore moves that look bad, and follow promising paths as far as possible (deep but narrow strategy)

# Questions?

# Cutoff Test

- It is still impractical to search down till terminal nodes and propagate utility values up

# Cutoff Test

- It is still impractical to search down till terminal nodes and propagate utility values up
- It is prudent to **cutoff** the search at some pre-determined depth (ply)  
[Either Type A or Type B strategy]

# Cutoff Test

- It is still impractical to search down till terminal nodes and propagate utility values up
- It is prudent to **cutoff** the search at some pre-determined depth (ply) [Either Type A or Type B strategy]
- A simple cutoff test replaces the terminal test — cutoff may be based on just the depth or other simple logic

# Cutoff Test

- It is still impractical to search down till terminal nodes and propagate utility values up
- It is prudent to **cutoff** the search at some pre-determined depth (ply) [Either Type A or Type B strategy]
- A simple cutoff test replaces the terminal test — cutoff may be based on just the depth or other simple logic
- But, what is the utility of a non-terminal cutoff node?

# Cutoff Test

- It is still impractical to search down till terminal nodes and propagate utility values up
- It is prudent to **cutoff** the search at some pre-determined depth (ply) [Either Type A or Type B strategy]
- A simple cutoff test replaces the terminal test — cutoff may be based on just the depth or other simple logic
- But, what is the utility of a non-terminal cutoff node? — use a **heuristic evaluation function**

# Cutoff Test

- It is still impractical to search down till terminal nodes and propagate utility values up
- It is prudent to **cutoff** the search at some pre-determined depth (ply) [Either Type A or Type B strategy]
- A simple cutoff test replaces the terminal test — cutoff may be based on just the depth or other simple logic
- But, what is the utility of a non-terminal cutoff node? — use a **heuristic evaluation function**

H-MINIMAX( $s, d$ ) =

$$\begin{cases} Eval(s, MAX) & \text{if IS-CUTOFF}(s, d) \\ \max_{a \in ACTIONS(s)} H-MM(RESULT(s, a), d + 1) & \text{if TO-MOVE}(s) = MAX \\ \min_{a \in ACTIONS(s)} H-MM(RESULT(s, a), d + 1) & \text{if TO-MOVE}(s) = MIN \end{cases}$$



# Heuristic Evaluation Function

- An evaluation function returns an estimate of the expected utility of a state  $s$  to player  $p$  — quality of this estimation is very important

# Heuristic Evaluation Function

- An evaluation function returns an estimate of the expected utility of a state  $s$  to player  $p$  — quality of this estimation is very important
- Should identify terminal states — order them the same way as the true utility function
- Value should be somewhere between a “loss” and a “win”
- Should be easy to compute!
- Strongly correlated with the actual chances of winning

# Heuristic Evaluation Function

- Evaluation functions may be designed based on certain **features** of a state

# Heuristic Evaluation Function

- Evaluation functions may be designed based on certain **features** of a state
- With features, we may abstract a set of states as a category (or equivalence classes) and estimate the expected value from experience

# Heuristic Evaluation Function

- Evaluation functions may be designed based on certain **features** of a state
- With features, we may abstract a set of states as a category (or equivalence classes) and estimate the expected value from experience
- For example, 82% of states encountered in the two-pawns vs one-pawn category lead to a win; 2% to a loss; and 16% to a draw — expected value may be  $(0.82 \times 1) + (0.02 \times 0) + (0.16 \times \frac{1}{2}) = 0.90$

# Heuristic Evaluation Function

- Evaluation functions may be designed based on certain **features** of a state
- With features, we may abstract a set of states as a category (or equivalence classes) and estimate the expected value from experience
- For example, 82% of states encountered in the two-pawns vs one-pawn category lead to a win; 2% to a loss; and 16% to a draw — expected value may be  $(0.82 \times 1) + (0.02 \times 0) + (0.16 \times \frac{1}{2}) = 0.90$
- Another simple idea: Consider a weighted combination of the features (each feature is assumed to contribute a numeric value)

# Heuristic Evaluation Function


- Evaluation functions may be designed based on certain **features** of a state
- With features, we may abstract a set of states as a category (or equivalence classes) and estimate the expected value from experience
- For example, 82% of states encountered in the two-pawns vs one-pawn category lead to a win; 2% to a loss; and 16% to a draw — expected value may be  $(0.82 \times 1) + (0.02 \times 0) + (0.16 \times \frac{1}{2}) = 0.90$
- Another simple idea: Consider a weighted combination of the features (each feature is assumed to contribute a numeric value)
- Neural networks basically compute weighted sum of the features and apply an non-linear activation function

# Heuristic Evaluation Function


- Evaluation functions may be designed based on certain **features** of a state
- With features, we may abstract a set of states as a category (or equivalence classes) and estimate the expected value from experience
- For example, 82% of states encountered in the two-pawns vs one-pawn category lead to a win; 2% to a loss; and 16% to a draw — expected value may be  $(0.82 \times 1) + (0.02 \times 0) + (0.16 \times \frac{1}{2}) = 0.90$
- Another simple idea: Consider a weighted combination of the features (each feature is assumed to contribute a numeric value)
- Neural networks basically compute weighted sum of the features and apply an non-linear activation function
- Neural networks and deep learning based chess engines are common today!



# Issues in cutoff




(a) White to move




(b) White to move

# Issues in cutoff




(a) White to move




(b) White to move

- Evaluation should be applied only to positions that are **quiescent**

# Issues in cutoff




(a) White to move




(b) White to move

- Evaluation should be applied only to positions that are **quiescent**
- **Quiescence search**: Perform extra search to confirm that there are no wild swings in the evaluation

# Issues in cutoff



# Issues in cutoff



- **Horizon effect**
- Keep a collection of **singular extensions** — allow moves that are “clearly better” than all other moves in a given position, even after cut-off

# Forward Pruning

- Alpha-Beta pruning is perfect — does not affect the completeness or optimality of the algorithm

# Forward Pruning

- Alpha-Beta pruning is perfect — does not affect the completeness or optimality of the algorithm
- It is also possible to think of imperfect pruning

# Forward Pruning

- Alpha-Beta pruning is perfect — does not affect the completeness or optimality of the algorithm
- It is also possible to think of imperfect pruning
- **Forward Pruning** considers only a few valid moves in a state (Type B Strategy!)



# Forward Pruning

- Alpha-Beta pruning is perfect — does not affect the completeness or optimality of the algorithm
- It is also possible to think of imperfect pruning
- **Forward Pruning** considers only a few valid moves in a state (Type B Strategy!)
- **Beam search** is an approach to forward pruning where only a “beam” of the  $n$  best moves (according to some heuristics) are considered

# Forward Pruning

- Alpha-Beta pruning is perfect — does not affect the completeness or optimality of the algorithm
- It is also possible to think of imperfect pruning
- **Forward Pruning** considers only a few valid moves in a state (Type B Strategy!)
- **Beam search** is an approach to forward pruning where only a “beam” of the  $n$  best moves (according to some heuristics) are considered
- **ProbCut** by Buro (1995) prunes nodes that are probably outside the current  $(\alpha, \beta)$  window (based on statistics gathered from experience)

# Forward Pruning

- Alpha-Beta pruning is perfect — does not affect the completeness or optimality of the algorithm
- It is also possible to think of imperfect pruning
- **Forward Pruning** considers only a few valid moves in a state (Type B Strategy!)
- **Beam search** is an approach to forward pruning where only a “beam” of the  $n$  best moves (according to some heuristics) are considered
- **ProbCut** by Buro (1995) prunes nodes that are probably outside the current  $(\alpha, \beta)$  window (based on statistics gathered from experience)
- **Late move reduction** — if the moves are ordered, probably moves that appear late in the sequence are not good, and so depth may be reduced for them

# Search Vs Lookup

- A judicious combination of all these techniques has proved to be useful in practice — good heuristic evaluation, cutoff with quiescence search, a large transposition table, learning from a large collection of game database

# Search Vs Lookup

- A judicious combination of all these techniques has proved to be useful in practice — good heuristic evaluation, cutoff with quiescence search, a large transposition table, learning from a large collection of game database — a 5-ply minimax search can be extended to about 14 ply search on a given hardware and time

# Search Vs Lookup

- A judicious combination of all these techniques has proved to be useful in practice — good heuristic evaluation, cutoff with quiescence search, a large transposition table, learning from a large collection of game database — a 5-ply minimax search can be extended to about 14 ply search on a given hardware and time
- For games like chess, **table lookup** for opening and end games have found to be very effective! — search techniques are required only for the middle game

# Search Vs Lookup

- A judicious combination of all these techniques has proved to be useful in practice — good heuristic evaluation, cutoff with quiescence search, a large transposition table, learning from a large collection of game database — a 5-ply minimax search can be extended to about 14 ply search on a given hardware and time
- For games like chess, **table lookup** for opening and end games have found to be very effective! — search techniques are required only for the middle game
- Reading Exercise: Study open source chess playing engines such as crafty and stockfish (<https://stockfishchess.org/>)

# Search Vs Lookup

- A judicious combination of all these techniques has proved to be useful in practice — good heuristic evaluation, cutoff with quiescence search, a large transposition table, learning from a large collection of game database — a 5-ply minimax search can be extended to about 14 ply search on a given hardware and time
- For games like chess, **table lookup** for opening and end games have found to be very effective! — search techniques are required only for the middle game
- Reading Exercise: Study open source chess playing engines such as crafty and stockfish (<https://stockfishchess.org/>)
- Reading Exercise: You may also optionally read about the recent developments based on neural networks and deep learning, such as chess engine Leela Chess Zero (<https://lczero.org/>)



# Questions?

# Monte Carlo Tree Search

- There are games where the branching factor is quite large and it is not easy to come up with an evaluation function — for example, Go

# Monte Carlo Tree Search

- There are games where the branching factor is quite large and it is not easy to come up with an evaluation function — for example, Go
- Go has a branching factor that starts with 361 and most of the states are in flux until the endgame

# Monte Carlo Tree Search

- There are games where the branching factor is quite large and it is not easy to come up with an evaluation function — for example, Go
- Go has a branching factor that starts with 361 and most of the states are in flux until the endgame
- So a different strategy called **Monte Carlo Tree Search (MCTS)** has been evolved

# Monte Carlo Tree Search

- MCTS evaluates a state  $s$  by conducting simulations of complete games (one line of play) starting from  $s$

# Monte Carlo Tree Search

- MCTS evaluates a state  $s$  by conducting simulations of complete games (one line of play) starting from  $s$
- Utility of  $s$  is estimated as the average utility over a number of simulations

# Monte Carlo Tree Search

- MCTS evaluates a state  $s$  by conducting simulations of complete games (one line of play) starting from  $s$
- Utility of  $s$  is estimated as the average utility over a number of simulations
- A simulation (also called as a playout or rollout) chooses moves for each player until a terminal position is reached (utility is known)

# Monte Carlo Tree Search

- MCTS evaluates a state  $s$  by conducting simulations of complete games (one line of play) starting from  $s$
- Utility of  $s$  is estimated as the average utility over a number of simulations
- A simulation (also called as a playout or rollout) chooses moves for each player until a terminal position is reached (utility is known)
- Utility of  $s$  will be average utility or win percentage (whichever makes sense) for all the simulations conducted so far



# Monte Carlo Tree Search

- MCTS evaluates a state  $s$  by conducting simulations of complete games (one line of play) starting from  $s$
- Utility of  $s$  is estimated as the average utility over a number of simulations
- A simulation (also called as a playout or rollout) chooses moves for each player until a terminal position is reached (utility is known)
- Utility of  $s$  will be average utility or win percentage (whichever makes sense) for all the simulations conducted so far
- For example, utility of  $s$  may look like  $\frac{22}{39}$ , which means MAX won the game 22 times among the 39 playouts

# Monte Carlo Tree Search


- MCTS evaluates a state  $s$  by conducting simulations of complete games (one line of play) starting from  $s$
- Utility of  $s$  is estimated as the average utility over a number of simulations
- A simulation (also called as a playout or rollout) chooses moves for each player until a terminal position is reached (utility is known)
- Utility of  $s$  will be average utility or win percentage (whichever makes sense) for all the simulations conducted so far
- For example, utility of  $s$  may look like  $\frac{22}{39}$ , which means MAX won the game 22 times among the 39 playouts
- If one more simulation is conducted from  $s$ , and MAX loses in that playout, then the utility of  $s$  is revised as  $\frac{22}{40}$

- A **playout policy** is used to select a good move for each player during a playout

- A **playout policy** is used to select a good move for each player during a playout
- When the game is in progress and a search tree is kept in memory, a **selection policy** is used to select a child at each step and decide where to start the playout from


- A **playout policy** is used to select a good move for each player during a playout
- When the game is in progress and a search tree is kept in memory, a **selection policy** is used to select a child at each step and decide where to start the playout from
- Selection policy may have two strategies: **Exploration** of states that have had few playouts, and **exploitation** of states that have done well in the past playouts

# MCTS: Selection




(a) Selection

# MCTS: Expansion and Simulation



## MCTS: Back-propagation



### (c) Backpropagation



```
function MONTE-CARLO-TREE-SEARCH(state) returns an action  
  tree  $\leftarrow$  NODE(state)  
  while IS-TIME-REMAINING() do  
    leaf  $\leftarrow$  SELECT(tree)  
    child  $\leftarrow$  EXPAND(leaf)  
    result  $\leftarrow$  SIMULATE(child)  
    BACK-PROPAGATE(result, child)  
  return the move in ACTIONS(state) whose node has highest number of playouts
```

# Upper confidence bound selection strategy

- We have seen that the selection policy needs to balance between exploitation and exploration

# Upper confidence bound selection strategy

- We have seen that the selection policy needs to balance between exploitation and exploration
- One way of achieving it will be through an **Upper confidence bound** formula called **UCB1**. For a node  $n$ , the formula is given by

$$UCB1(n) = \frac{U(n)}{N(n)} + C \times \sqrt{\frac{\log N(Parent(n))}{N(n)}}$$

# Upper confidence bound selection strategy

- We have seen that the selection policy needs to balance between exploitation and exploration
- One way of achieving it will be through an **Upper confidence bound** formula called **UCB1**. For a node  $n$ , the formula is given by

$$UCB1(n) = \frac{U(n)}{N(n)} + C \times \sqrt{\frac{\log N(Parent(n))}{N(n)}}$$

- The value of constant  $C$  is usually taken as  $\sqrt{2}$

# Upper confidence bound selection strategy

- We have seen that the selection policy needs to balance between exploitation and exploration
- One way of achieving it will be through an **Upper confidence bound** formula called **UCB1**. For a node  $n$ , the formula is given by

$$UCB1(n) = \frac{U(n)}{N(n)} + C \times \sqrt{\frac{\log N(Parent(n))}{N(n)}}$$

- The value of constant  $C$  is usually taken as  $\sqrt{2}$
- It is possible to use **early playout termination**, where a non-terminal is evaluated by a heuristic function

- It is possible to use MCTS for new games, in which there is no body of experience or database of games
- Knowledge of the game can be encoded into selection and playout policies

- It is possible to use MCTS for new games, in which there is no body of experience or database of games
- Knowledge of the game can be encoded into selection and playout policies
- MCTS predominantly uses Type B strategy and so may completely miss a good line of play
- Even “obvious” states may need several playouts to estimate the utility

- It is possible to use MCTS for new games, in which there is no body of experience or database of games
- Knowledge of the game can be encoded into selection and playout policies
- MCTS predominantly uses Type B strategy and so may completely miss a good line of play
- Even “obvious” states may need several playouts to estimate the utility
- MCTS is a kind of reinforcement learning



# Questions?