

AI - Last Min Notes

Owner	N nithu
Tags	

Unit 1 - AGENTS AND BASIC SEARCH TECHNIQUES

AGENTS

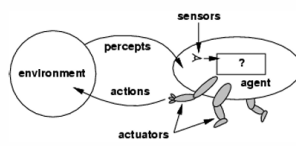
A rational agent → act to achieve the best outcome, when there is uncertainty.

Should learn from percepts and exploration

Autonomous

Agent- perceives env. thro sensors and acts thro actuators

Eg: Human Body



Consequentialism - evaluate an agent's behaviour thro consequences.

Perf. measure - criteria for success of an agent

P- Performance Measure

E- environment

A- Actuators

S-Sensors

PROPERTIES OF ENVIRONMENT

Fully observable vs partially observable(Chess vs Kriegspiel)

Single vs Multi

Deterministic vs non-deterministic(stochastic) - determine the next state with the current state vs multiple unpredictable outcomes

Episodic vs sequential - action taken once is dependent on previous

Static vs Dynamic - changing the env.

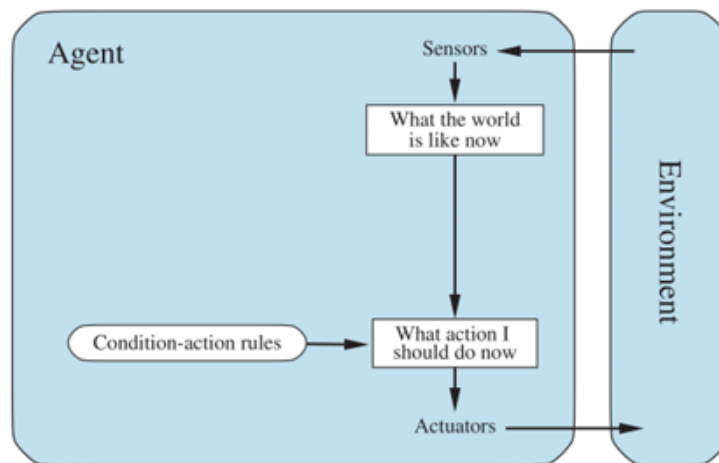
Discrete vs cont: time is discrete or cont.

Known or unknown

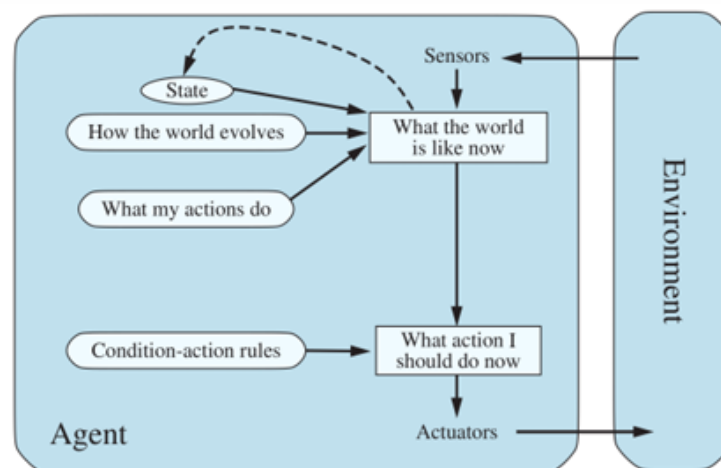
TYPES OF AGENTS:

Simple reflex(Do not have mem. of past)→ completely dependent on present.

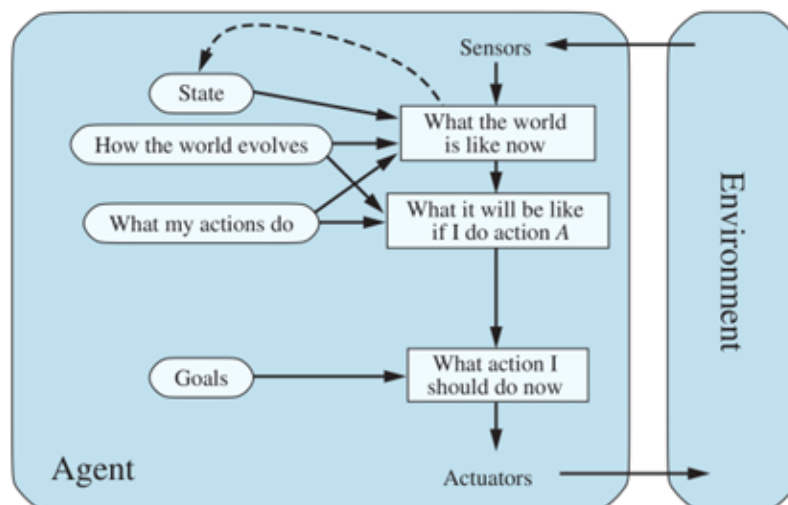
Condition-action rules



Model based reflex(Internal state → past states, work on change)



Goal based agents/Problem solving agents(Desirable goal to be achieved → future events and results)



Steps for goal-based agents:

1. Goal Formulation: Set of one/more world states
2. Prob. Formulation: What to consider and what not to consider
3. Search for a solution - a sequence of actions from the initial to the goal state.
4. Execute

→ The process of looking for the best soln → search

→ the best seq(list of actions) → soln

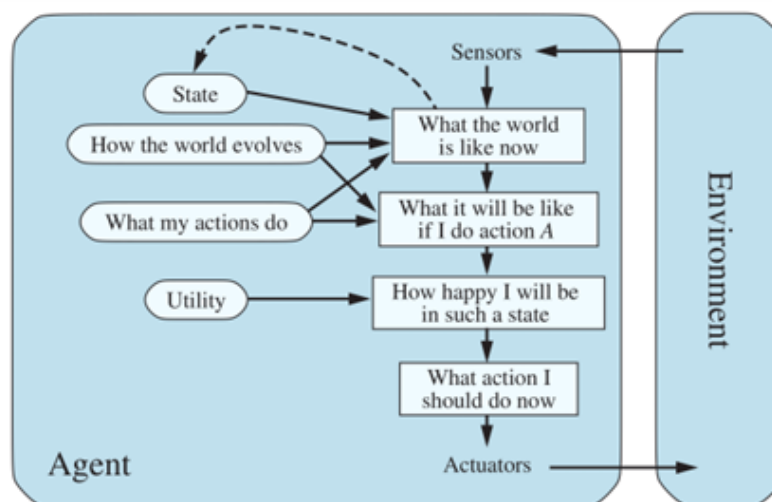
TYPES OF PROBLEMS

1. Deterministic, fully observable
2. Non-observable
3. Non-Dete. and partially observable
4. Unknown state space

1. Initial state
2. Actions
3. Transition model or (Successor functions)
4. Goal Test.
5. Path Cost.

Optimal soln → lowest path cost among the lot

Utility based agents (Alternatives → indicating a measure of success or happiness)



Learning agents (Reinforcement based)

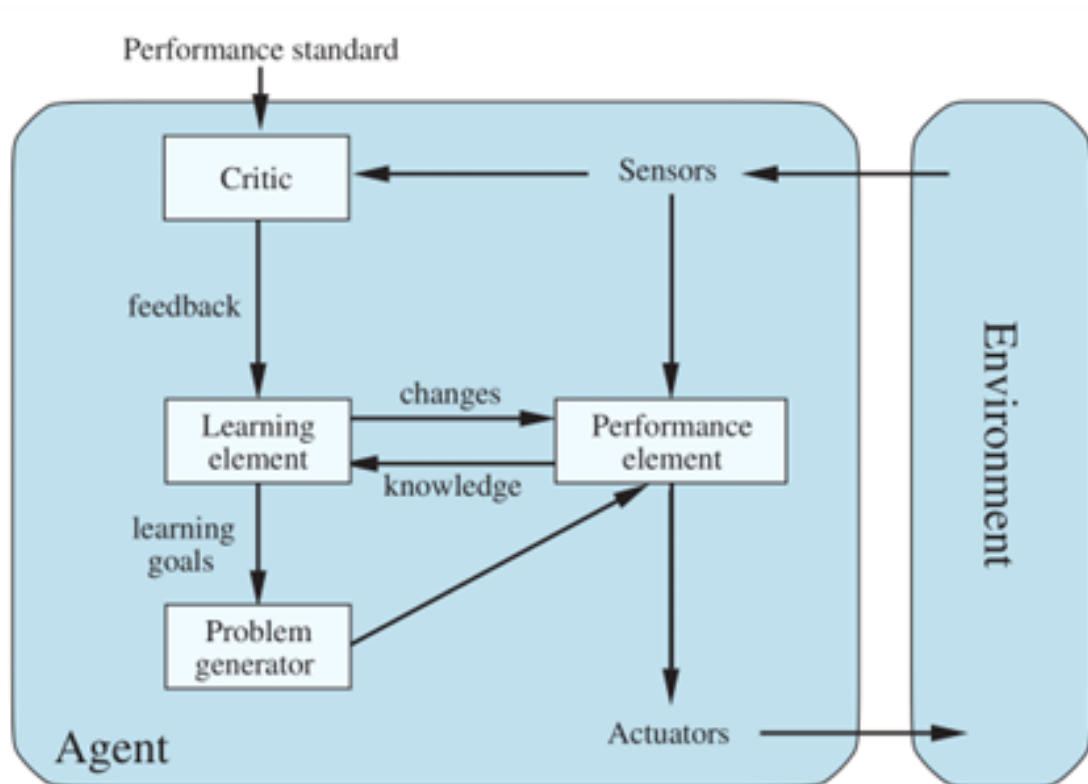


Table-lookup driven agent —: implements a lookup table for actions. Too large
 → 4^K possible seq.

State-Space Approach: Visualize the search tree based on an initial state, set of actions

B.F → no of possible actions

State Space → all that is reachable from the initial state

State-space components: Initial state, Successor function, Goal test and Path cost

8-puzzle prob(jumbled up numbers from 1-8 → arrange it ascendingly)

- Uninformed:
 - Breadth-First
 - Depth-First
 - Depth-Limited
 - Bi-Directional Search
- Informed (Heuristics):
 - Best-first Greedy
 - A^*
 - Local Search Strategies
- Constraint Satisfaction

Metrics of a problem:

Completeness

Time Complexity

Space Complexity

Optimality

UNINFORMED SEARCH TECH.

Tree-Search -generate successors of the explored states

Graph-search - explored and frontier(generated but not explored)

States (physical config.)vs Nodes (state, parent node, action,path cost,depth)

BFS →FIFO Queue

Complete: Yes(if B.F is finite)

Time Complexity: $O(b^{(d+1)})$

Space Complexity: $O(b^{(d+1)})$

Optimality: No

Uniform cost search → Start with the lowest cost fringe

Complete: Yes(if no negative cost)

Time Complexity: $O(b^{(C^*/e)})$

Space Complexity: $O(b^{(C^*/e)})$

Optimality: Yes. Monotonic cost fn.

[Use I.D idea within Uniform cost search - Iterative Lengthening Approach]

Depth-First Search → LIFO Stack

Complete: No

Time Complexity: $O(b^m)$

Space Complexity: $O(b(m))$

Optimality: No

Backtracking → DFS version → generate one successor

Complete: No

Time Complexity: $O(b^m)$

Space Complexity: $O(m)$

Optimality: No

What we need now?

Something with less time complexity as DFS and complete as BFS →

Depth-Limited Search : Impose a limit l on the search

$l > d$ but not large (defeats the purpose)

Complete: No! Depends on L

Optimal: No!

Time Complexity = $O(b^l)$

Space Complexity: $O(bl)$

Iterative Deepening: Dynamically find a suitable depth-limit l and increment until goal is reached

Complete: Yes! - BFS Property

Optimal: No! - BFS Property

Time Complexity = $O(b^d)$ - DFS Property

Space Complexity: $O(bd)$ - DFS Property

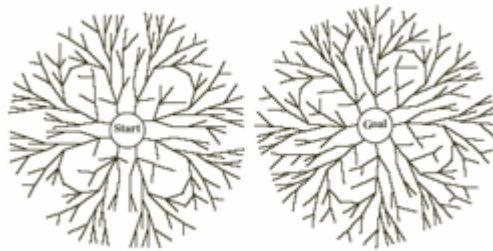
Bidirectional search: Run two search simultaneously - one forward and another backward from the goal - make them meet in middle

Time Complexity = $O(b^{d/2})$ - half the depth

Space Complexity = $O(b^{d/2})$

Goal test → 'intersection test'

Backward search → define predecessors



UNIT 2 - HEURISTICS, INFORMED SEARCH AND LOCAL SEARCH

HEURISTICS - A* search/RBFS/BEST-FIRST - INFORMED SEARCH

BEST-first GREEDY ALGORITHM:

evaluation function=heuristics

Complete: No

Optimal: No

Time Complexity: $O(b^m)$

Space Complexity: $O(b^m)$

What if we design an evaluation fn - to determine goodness or badness of a state. → to choose the best of the lot

→ Need of heuristics $h(s)$ → estimated cost of reaching a goal state

An admissible heuristic → never overestimates the cost → optimistic

A heuristic h is admissible if for every node n , $h(n) \leq h^*(n)$, where $h^*(n)$ is the true cost to reach the goal state from n .

A heuristic is consistent if for every node n , every successor n' of n generated by any action a , means it is non-decreasing along a path. f is monotonic

DOMINANT HEURISTICS: If $h_2(n) \geq h_1(n)$ then h_2 is said to be dominant, which makes it a better heuristic for search.

This helps in reducing the search-space by suitable generating successors, with a good branching factor.

How to get heuristics?

- Relaxed problems:
- Subproblems: like in 4-queens problems. Use the substeps to minimize the heuristic value
- Landmarks: Optimal paths between two nodes using precomputation. Use landmark points among the vertices.

Differential heuristic:

$$h_{DH}(n) = \max_{L \in \text{Landmarks}} |C^*(n, L) - C^*(goal, L)|$$

A* search:

- Optimal
- Complete
- Time complexity: $O(b^d)$
- Optimally efficient(the fewest nodes are expanded)
- Space Complexity \rightarrow worst case : BFS= $O(b^d)$
- If admissible heuristic - use A* tree-search
- If consistent heuristic - use A* graph-search

BEAM SEARCH : Limiting the frontier

Have only the k-best nodes with f-scores and discard the others or have a limit for the f-scores

ITERATIVE - DEEPENING A* \rightarrow f-cost is a cutoff(like iterative lengthening)—
>path-finding

RBFS : Dynamically adjusting cost-limit for a path

Track the best alternative path , when the best child exceeds, unwind

Better than IDA* , but still suffers from node regeneration.

SMA(Simplified Memory-bounded A*) \rightarrow Memory use to the peak

Proceed like A* until mem is exhausted

Throw away the node with the largest f-val to get new space.

Remember the best forgotten ones in the parent

Same f-cost is an issue here

LOCAL SEARCH TECHNIQUES

Our goal state is our final aim- so sometimes it is better to move on - iterating upon optimal config. satisfying a few set conditions rather than discussing your path.

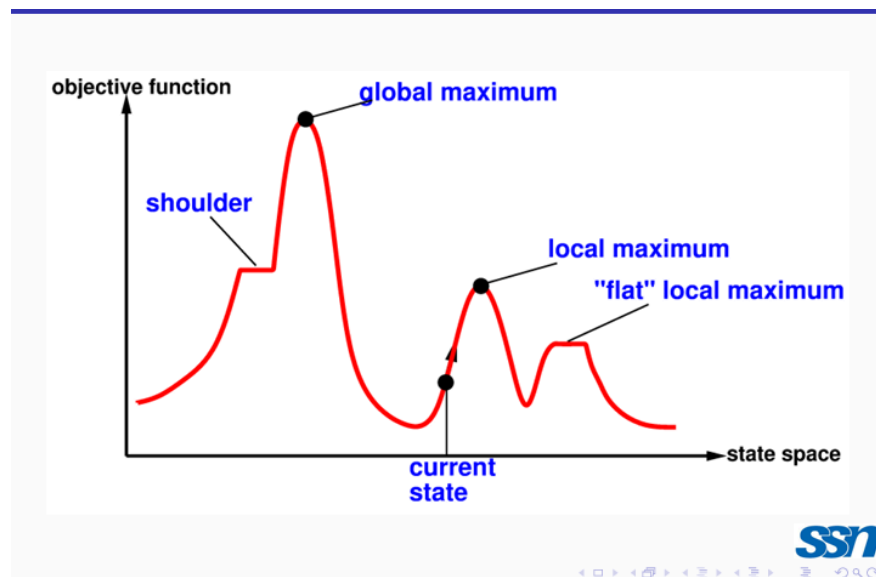
Eg: n-queens

1. Start with a state(any random placement of the n-queens)
2. Place one queen per column
3. Calculate and recalculate your heuristic
4. Iterate until the val comes at a minimum.

→ Iterative Improvement Algos.

→ Basically improve upon what you have in your hand

Hill climbing:



Problems:

1. Local Maxima/Minima - won't move further even if a better solution is available(kind of like instant gratification)
2. Ridges - Sequence of local maxima - difficult to navigate
3. Plateau - a very flat sequence or a shoulder

How to rectify?

Random sideways moves - you could escape a shoulder, but gets stuck on the maxima.

Stochastic Hill Climbing: Choose random out of all possibilities(prob. depends on steepness)

First Choice Hill Climbing: randomly generate your successors, until the immediate better option is generated

Random Restarts: Expect $1/p$ restarts for p prob. of success

Stimulated Annealing: A variation of hill climbing

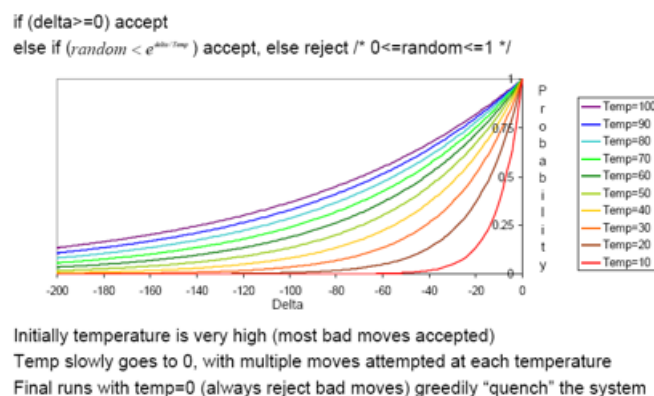
```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  current  $\leftarrow$  problem.INITIAL
  for  $t = 1$  to  $\infty$  do
     $T \leftarrow$  schedule( $t$ )
    if  $T = 0$  then return current
    next  $\leftarrow$  a randomly selected successor of current
     $\Delta E \leftarrow$  VALUE(current) - VALUE(next)
    if  $\Delta E > 0$  then current  $\leftarrow$  next
    else current  $\leftarrow$  next only with probability  $e^{-\Delta E/T}$ 
```

Start with a case where the prob. of a bad state is high and iterate to slowly reduce it.

Tries to explore enough early on, makes certain downhill moves for finding a good way uphill.

The prob. of moving to a higher energy state

$$p = e^{(-\Delta E/kT)}$$



LOCAL BEAM SEARCH: Generate k initial states terminate if goal state else generate all successors and select best k

GENETIC ALGORITHM: state with a random pop. Produce a new gen by choosing crossover pts and mutation

Schema : a substring where some positions are left undefined.

-
- Given: population **P** and fitness-function **f**
 - repeat
 - **newP** \leftarrow empty set
 - for $i = 1$ to **size(P)**
 - $x \leftarrow \text{RandomSelection}(\mathbf{P}, \mathbf{f})$
 - $y \leftarrow \text{RandomSelection}(\mathbf{P}, \mathbf{f})$
 - $child \leftarrow \text{Reproduce}(x, y)$
 - if (small random probability) then $child \leftarrow \text{Mutate}(child)$
 - add $child$ to **newP**
 - **P** \leftarrow **newP**

NON-DETERMINISM, PARTIALLY OBSERVABLE AND CONTINUOUS SPACE

In case of continuous state spaces,

Choose hill climbing or stimulated annealing because all the successors need not be generated.

The evaluation function $f(s)$ may be easily computed as:

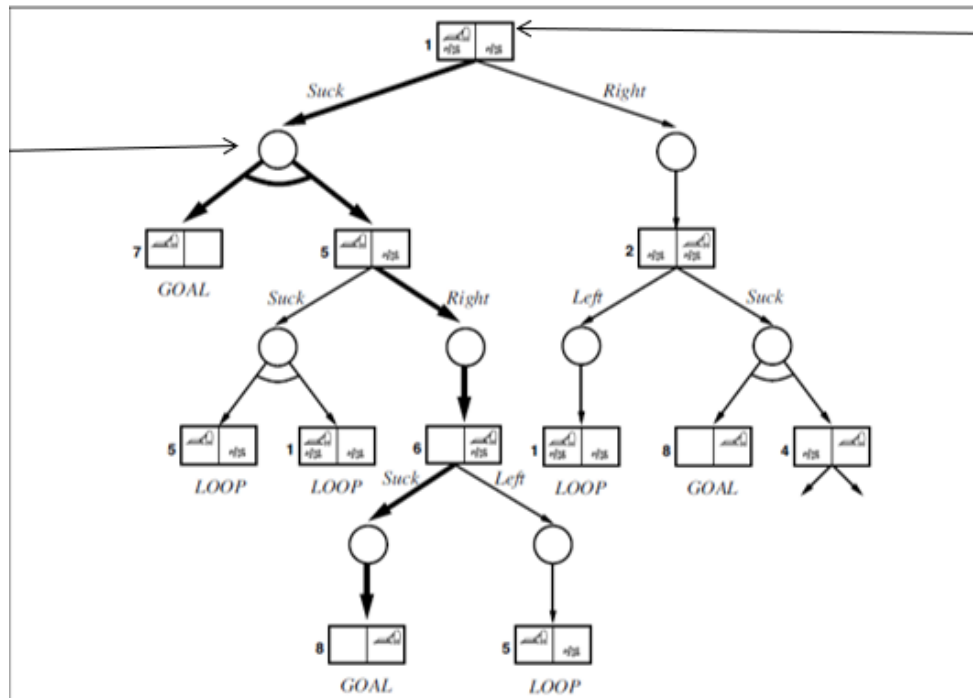
$$f((x_1, y_1), (x_2, y_2), (x_3, y_3)) = \sum_{i=1}^3 \sum_{c \in C_i} (x_i - x_c)^2 + (y_i - y_c)^2$$

Use empirical gradient /gradient descent to find the max/min in a landscape. Gives the steepest slope.

or take up Newton-Raphson method

The results of an action can vary in such a case . For example, suppose a vacuum cleaner cleans two squares of dirt, or deposits dirt while cleaning

Generalize the solution to a contingency plan. AND-OR TREE



In case of an partially observable environment(for eg, Kriegspiel) → the agent does not always know its state, instead goes by a belief state

ONLINE SEARCH AGENTS:

Executes actions

Acquires percepts

Deterministic and fully observable

Has to act to know the outcome

Eg: Autonomous vehicle and Search engines

```

function ONLINE-DFS-AGENT(problem, s') returns an action
    s, a, the previous state and action, initially null
    persistent: result, a table mapping (s, a) to s', initially empty
    untried, a table mapping s to a list of untried actions
    unbacktracked, a table mapping s to a list of states never backtracked to

    if problem.IS-GOAL(s') then return stop
    if s' is a new state (not in untried) then untried[s'] ← problem.ACTIONS(s')
    if s is not null then
        result[s, a] ← s'
        add s to the front of unbacktracked[s']
    if untried[s'] is empty then
        if unbacktracked[s'] is empty then return stop
        else a ← an action b such that result[s', b] = POP(unbacktracked[s'])
    else a ← POP(untried[s'])
    s ← s'
    return a
  
```

UNIT -3 GAME THEORY - ALPHA-BETA PRUNING, MCTS, CSP

All the search techniques that we have discussed so far , will seem fruitless for a actual game

GAMES:

Perfect info+deterministic → Chess

Imperfect info + deterministic → Kriegspiel

Perfect info + non-deterministic → Backgammon

Imperfect info + non-deterministic → Poker, scrabble

ASSUMPTIONS:

- Fully observable
- Adversarial
- Agents whose actions are alternate

For eg: 35^{100} possibilities for a move(non-trivial games)

Search → Path(not adversarial)

Game → a strategy(adversarial → your opp. plans as you do)

ZERO-SUM GAMES: two outcomes , one maximum/positive(+1) and one negative/minimum(-1) → Tic-tac toe

Reward+punishment(penalty) is always a constant

- States and the **initial state** S_0
- **TO_MOVE(s)**: The player whose turn it is to move in state s
- **ACTIONS(s)**: set of legal moves in a state s
- **RESULT(s,a)**: the **transition model** that defines the result of a move a in a state s
- **IS_TERMINAL(s)**: Terminal test — 'true' when the game is over — terminal states
- **UTILITY(s,p)**: Utility function — an objective function that defines the final numeric value to a player p when the game ends in a terminal state s — in chess, outcome is a win, loss, or draw, with values $+1, 0, \frac{1}{2}$

MINIMAX TREE(Adversarial Search) + PRUNING

MINIMAX Tree → alternate MIN and MAX (choose lowest utility for MIN and highest utility for MAX)

Heuristic Evaluation Fn: How 'good' a config. is for a player

```
function MINIMAX-SEARCH(game, state) returns an action
  player ← game.TO-MOVE(state)
  value, move ← MAX-VALUE(game, state)
  return move

function MAX-VALUE(game, state) returns a (utility, move) pair
  if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
  v ← -∞
  for each a in game.ACTIONS(state) do
    v2, a2 ← MIN-VALUE(game, game.RESULT(state, a))
    if v2 > v then
      v, move ← v2, a
  return v, move

function MIN-VALUE(game, state) returns a (utility, move) pair
  if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
  v ← +∞
  for each a in game.ACTIONS(state) do
    v2, a2 ← MAX-VALUE(game, game.RESULT(state, a))
    if v2 < v then
      v, move ← v2, a
  return v, move
```

Complete (if the tree is finite)

Optimal - if the opponent is optimal (he/she takes the best move)

Space Complexity: $O(bm)$

Time Complexity: $O(b^m)$

PRUNING: Discard based on a given heuristic.

ALPHA-BETA PRUNING:

- No use expanding bad nodes
- If the opp forces a bad node, then again its useless

ALPHA (Best choice for MAX) → lower bound for MAX

BETA (best choice for MIN) → upper bound for MAX

PRUNING CONDITION: $\alpha \geq \beta$

```

function ALPHA-BETA-SEARCH(game, state) returns an action
  player ← game.TO-MOVE(state)
  value, move ← MAX-VALUE(game, state, -∞, +∞)
  return move

function MAX-VALUE(game, state, α, β) returns a (utility, move) pair
  if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
  v ←  $-\infty$ 
  for each a in game.ACTIONS(state) do
    v2, a2 ← MIN-VALUE(game, game.RESULT(state, a), α, β)
    if v2 > v then
      v, move ← v2, a
      α ← MAX(α, v)
    if v ≥ β then return v, move
  return v, move

function MIN-VALUE(game, state, α, β) returns a (utility, move) pair
  if game.IS-TERMINAL(state) then return game.UTILITY(state, player), null
  v ←  $+\infty$ 
  for each a in game.ACTIONS(state) do
    v2, a2 ← MAX-VALUE(game, game.RESULT(state, a), α, β)
    if v2 < v then
      v, move ← v2, a
      β ← MIN(β, v)
    if v ≤ α then return v, move
  return v, move

```

Implementing this retains your completeness ,optimality and space complexity

Time complexity → can be halved $O(b^{m/2}) \rightarrow O(\sqrt{b}^m)$

Eg: In chess → B.F goes from 35 to $\sqrt{35}$ → approx 6

Worst-case → ordered branches →no pruning

Best Case → the left-most child is the best move(hence it is considered first)

HOW TO MAKE THIS EFFECTIVE?

Dynamic move ordering — moves found to be best in the past
(learning) — or, iterative-deepening

Transposition table — hash the α - β values for future use

Sometimes, this may also not be useful → Cases like Chess and Go

TYPE A → heuristic evaluation in a certain depth and estimate utilities

TYPE B → ignore bad moves

Iterative Deepening Search(IDS)

- **In practice, iterative deepening search (IDS) is used**

–IDS runs depth-first search with an increasing depth-limit

–when the clock runs out we use the solution found at the previous depth limit

NON DETERMINISM IN GAMES - MCTS

Basically - u could have multiple next states and work upon belief states or possible states

Monte Carlo Tree Search:

Four basic principles:

- Selection

Leaf node is a node which has unexplored child nodes.

$$UCB(node_i) = \bar{x}_i + c\sqrt{\frac{\log N}{n_i}}$$

$U(n)/C(n)$, $n \rightarrow$ node value , $N \rightarrow$ parent value of the node

As no of turns go up, the second term becomes lesser \rightarrow least likely to be selected

Built-in exploration and exploitation

- Expansion

Randomly pick an unexplored node

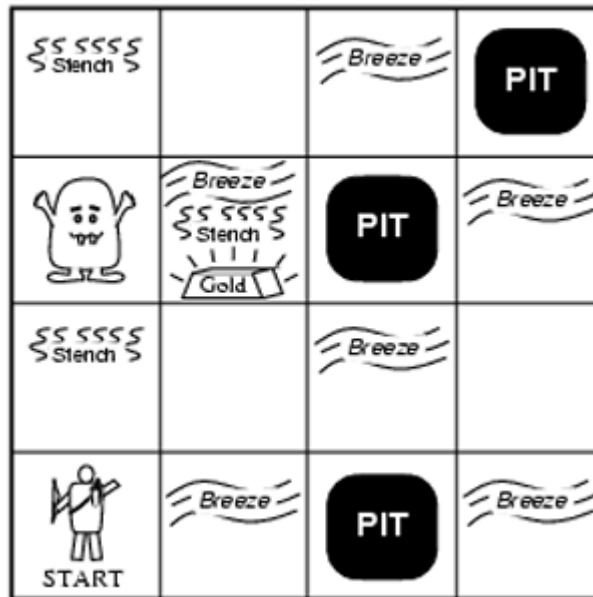
- Simulation

Rollout $\rightarrow +1/0/-1$

- Backpropagation

This can be considered as a Reinforcement Strategy \rightarrow using Type B strategy

WUMPUS WORLD



Breeze near Pit

Stench near Wumpus

Gold la Glitter

- Partially Observable
- Deterministic
- Episodic(No)
- Static
- Discrete
- Single Agent

SAT → to check for satisfiability → exists an assignment of truth values where the proposition holds true(NP-complete)

CSP

How does it differ from Local Search/Adversarial techniques?

→ Rather than consider a node is a goal state alone, we deep dive into the structure of the nodes and exploit that

A solution to the CSP → assignment of values to the variable such that all constraints are satisfied. Inconsistent, if nothing as such exists.

- A *state* is an *assignment* of values to some or all variables.

–An assignment is *complete* when every variable has a value.

An assignment is *partial* when some variables have no values

For example: For n-queens , the constraints are , $Q_i \neq Q_j$ and $|Q_i - Q_j| \neq |i - j|$

Variables: The regions

CSP - For graph coloring

Domain: The colors

Constraint: Adj.regions must have diff. colors

Types of Constraints:

Unary: $X = \text{red}$

Binary: $X = Y$

Higher order constraint: 3 or more variables

Preferences: Soft constraints

CSP → Can be expressed as a search problem

CSPs are commutative(the order doesnt matter)

DFS in CSP → backtracks

implement a constraint, backtracks when a variables has no legal values to assign

Algo:

SELECT UNASSIGNED VARIABLE

Min.Rem Values → most constrained variables and most likely to fail soon

If a tie, choose the var involved in the most constraints

ORDER-DOMAIN-VALUES

Least Constraining Value

Rules out the fewest choices for the variables it is in constraints with

FORWARD CHECKING: Take a constraint, change the variables → move on similarly

Sometimes, this can ignore inconsistencies

Arc Consistency

Two things are joined by an arc,all the variables in the two entities are consistent



UNIT 4 - PREDICATE LOGIC

Horn Clause \rightarrow At most positive literals

Definite clause \rightarrow Exactly one positive literal

Goal clause \rightarrow No positive literals

Forward chaining \rightarrow Bottom to top

Backward chaining \rightarrow Top to bottom

TABLE 1 Rules of Inference.

Rule of Inference	Tautology	Name
$\begin{array}{l} p \\ p \rightarrow q \\ \hline \therefore q \end{array}$	$(p \wedge (p \rightarrow q)) \rightarrow q$	Modus ponens
$\begin{array}{l} \neg q \\ p \rightarrow q \\ \hline \therefore \neg p \end{array}$	$(\neg q \wedge (p \rightarrow q)) \rightarrow \neg p$	Modus tollens
$\begin{array}{l} p \rightarrow q \\ q \rightarrow r \\ \hline \therefore p \rightarrow r \end{array}$	$((p \rightarrow q) \wedge (q \rightarrow r)) \rightarrow (p \rightarrow r)$	Hypothetical syllogism
$\begin{array}{l} p \vee q \\ \neg p \\ \hline \therefore q \end{array}$	$((p \vee q) \wedge \neg p) \rightarrow q$	Disjunctive syllogism
$\begin{array}{l} p \\ \hline \therefore p \vee q \end{array}$	$p \rightarrow (p \vee q)$	Addition
$\begin{array}{l} p \wedge q \\ \hline \therefore p \end{array}$	$(p \wedge q) \rightarrow p$	Simplification
$\begin{array}{l} p \\ q \\ \hline \therefore p \wedge q \end{array}$	$((p) \wedge (q)) \rightarrow (p \wedge q)$	Conjunction
$\begin{array}{l} p \vee q \\ \neg p \vee r \\ \hline \therefore q \vee r \end{array}$	$((p \vee q) \wedge (\neg p \vee r)) \rightarrow (q \vee r)$	Resolution

- Negation or De Morgan's Law (we saw this last time):
 - $\neg \forall x P(x) \equiv \exists x \neg P(x)$
 - $\neg \exists x P(x) \equiv \forall x \neg P(x)$
- Distributivity:
 - $\forall x (P(x) \wedge Q(x)) \equiv \forall x P(x) \wedge \forall x Q(x)$
 - $\exists x (P(x) \vee Q(x)) \equiv \exists x P(x) \vee \exists x Q(x)$
 - Can't distribute universal quantifier over disjunction or

UNIT 5 - UNCERTAINTY, BN

Uncertainty → partial observability, noisy sensors, uncertainty in actions and immense complexity in modelling

PROBABILITY: Model agent's degree of belief

Relates to agent's own state of knowledge, with obviously change with new evidence.

Utility Theory → represent preferences

Decision Theory = probability theory + utility theory

Atomic event: A complete specification of the state of the world.

- Mutually exclusive
- the set of all possible atomic events is exhaustive
- Entails the truth or falsehood of every proposition
- Logically equivalent to the disjunction of all atomic events

Prior probability: Or conditional probability

$$P(a \mid b) = P(a \wedge b) / P(b) \text{ if } P(b) > 0$$

$$P(a \wedge b) = P(a \mid b) P(b) = P(b \mid a) P(a)$$

Joint prob.dist: Prob. of Every atomic event of the random variables

Kolmogorov's Axioms:

$$0 \leq P(A) \leq 1$$

$$P(\text{true}) = 1 \text{ and } P(\text{false}) = 0$$

$$P(A \vee B) = P(A) + P(B) - P(A \wedge B)$$

$$P(\sim A) = 1 - P(A)$$

Chain rule = Product rule applied multiple times

Marginal Probability/Conditional Probability= $P(Y)=P(Y/z)*P(z)$

Time complexity: $O(d^m)$

Space Complexity: $O(d^n)$

Independence: $P(A|B)=P(A)$ or $P(B)$

$P(A,B)=P(A)P(B)$

Bayes' rule: Posterior= $P(a | b) = P(b | a) P(a) / P(b)$

Naive Bayes Model= $P(Y=1|X)=P(y=1)P(x_1|y=1).../P(X)$

Bayes Classifier: $\text{argmax} P(Y|X_1...X_n)$

The no of parameters for modelling go down significantly

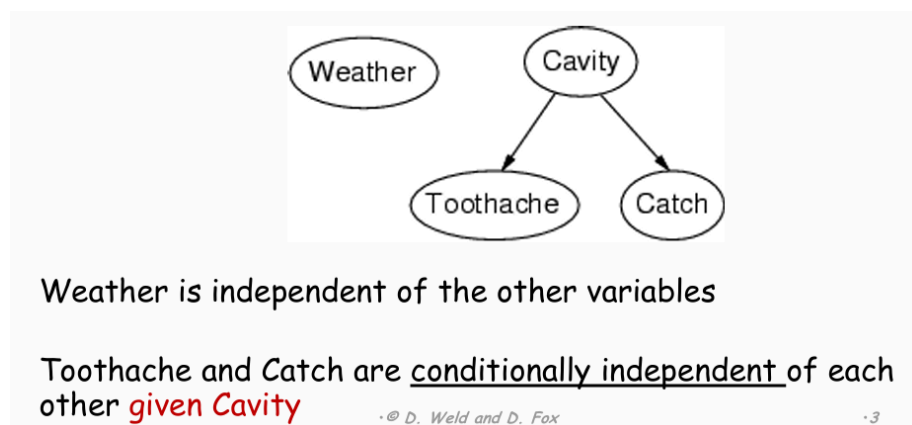
Applications: Spam classification, Weather, Medical diagnosis

→ Assesses diagnostic probability from casual probability

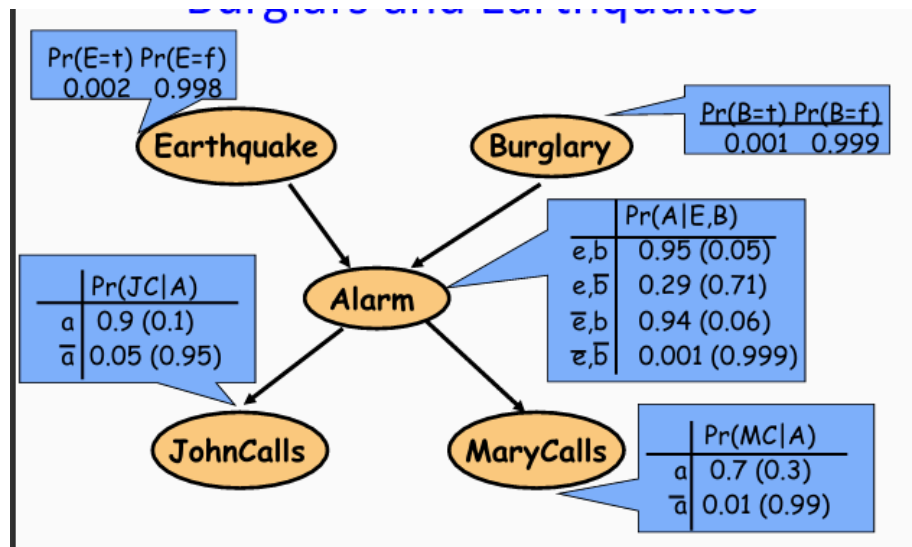
Bayes Nets(BN)

→ Join dist P over a set of variables requires exponential space for representation and inference

Graphical representation off conditional independence relations → BNs



→ It is basically a directed , acylic graph, with nodes, one per R.V



$$P(JC, MC, A, E, B) = P(JC|A) \cdot P(MC|A) \cdot P(A|B, E) \cdot P(E) \cdot P(B)$$

Full Joint Dist(Chain Rule)

$$P(X_1, X_2, \dots, X_n) = \prod_{i=1}^n P(X_i | \text{Par}(X_i))$$

Edges denote probabilistic influence

Marginal Prob from BN: $\Pr(Z|E) \rightarrow Z$: Query variable

E: Evidence variable

everything else: hidden

Variable Elimination: Eliminate everything until there is only the query variable

→ join all the factors

→ sum out the influence

→ exploit the product form

Inference in a BN is NP-hard (3-SAT into BN)

Probability = $1/2^n$

→ Polynomial time and compare with inference in CSPs

BN IN A GENERATIVE MODEL

Generate one variable in topological order

Stochastic Simulation → one after other

$P(b|c) = \# \text{ of live samples with } B=b / \text{total \# of live samples}$

Rejection Sampling → Returns only consistent posterior estimates

Likelihood weighting → fix evidence variables, sample only non-evidence, weigh each sample by likelihood of evidence

$S(z,e) = \prod P(z \setminus \text{Parents}(Z))$

$w(z,e) = \prod P(e \setminus \text{Parents}(E))$

Weighted Sampling Prob = $S(z,e)w(z,e)$

$= P(z,e)$

$$P(b|c) = \frac{\text{weight of samples with } B=b}{\text{total weight of samples}}$$

MCMC with Gibbs Sampling

- Fix the observed variables
- Randomly assign values to the non-observed
- Random walk → Pick X , Calc $P(X=\text{true} \setminus \text{other variables})$, Set $X=\text{true}$

Basically cal $P(x | \text{MB}(X))$

MB → Markov Blanket

A variable is independent of all others given its Markov Blanket

$$P(X) = \alpha P(X | \text{Parents}(X)) \prod_{Y \in \text{Children}(X)} P(Y | \text{Parents}(Y))$$

$$P(X|E) = \frac{\text{number of samples with } X=x}{\text{total number of samples}}$$

- **Advantages:**

- No samples are discarded
- No problem with samples of low weight
- Can be implemented very efficiently
 - 10K samples @ second

- **Disadvantages:**

- Can get stuck if relationship between two variables is *deterministic*
- Many variations have been devised to make MCMC more robust